

Apache Kafka Basics

David Veres

Motivation

- ▶ Event-driven systems are becoming more popular
- ▶ Focus on data as events
- ▶ Producers and consumers are decoupled
- ▶ Process events, react to them
- ▶ Needs a single platform
 - ▶ Connects the participants
 - ▶ Real-time
 - ▶ Stores the events

Apache Kafka



PayPal




Hotels.com™



NETFLIX

Fundamentals

Events

- ▶ „A thing that has happened”
- ▶ It can be any kind of thing
 - ▶ A user clicks on a button
 - ▶ A microservice finishes a computation and submits the result
 - ▶ An IoT device sends data
- ▶ An event has a state
 - ▶ Describes what happened
 - ▶ Stored in structured format (like JSON)
- ▶ In Kafka it is a key-value pair

Event
Key
Value

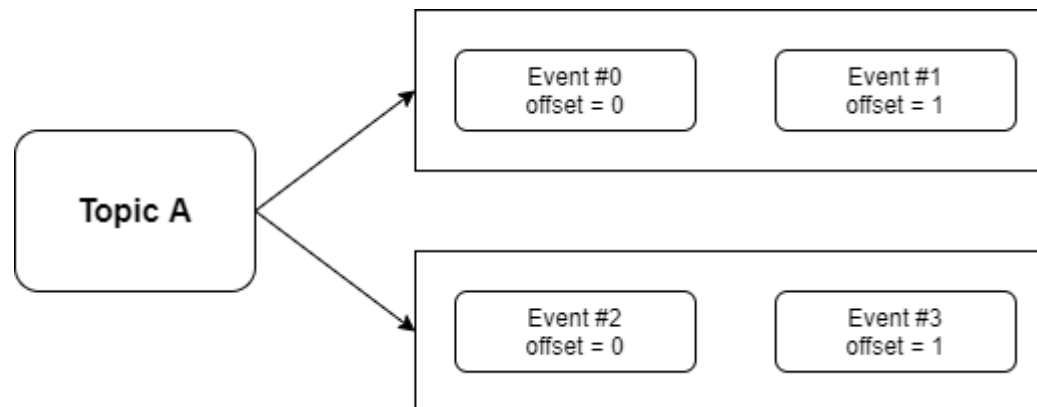
Topics

- ▶ The way of categorizing events
- ▶ A collection of similar events
- ▶ We can have multiple topics
- ▶ The same event can appear in different topics
- ▶ Topic = log of events
 - ▶ Append only
 - ▶ Events are immutable




Partitions

- ▶ A topic can be split into multiple partitions
- ▶ „Real log of events”
 - ▶ The order of events is strict
 - ▶ This might not true for topics
- ▶ Every partition has it's own offset space



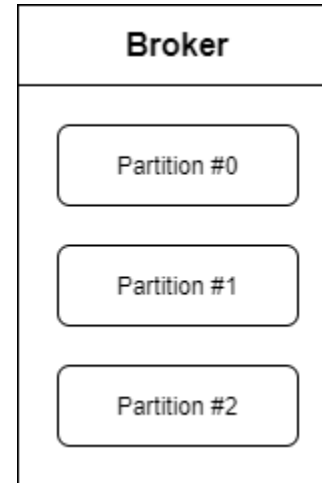
Segments

- ▶ A partition is broken into multiple segments
- ▶ These are individual files
- ▶ Events are stored on the disk in files



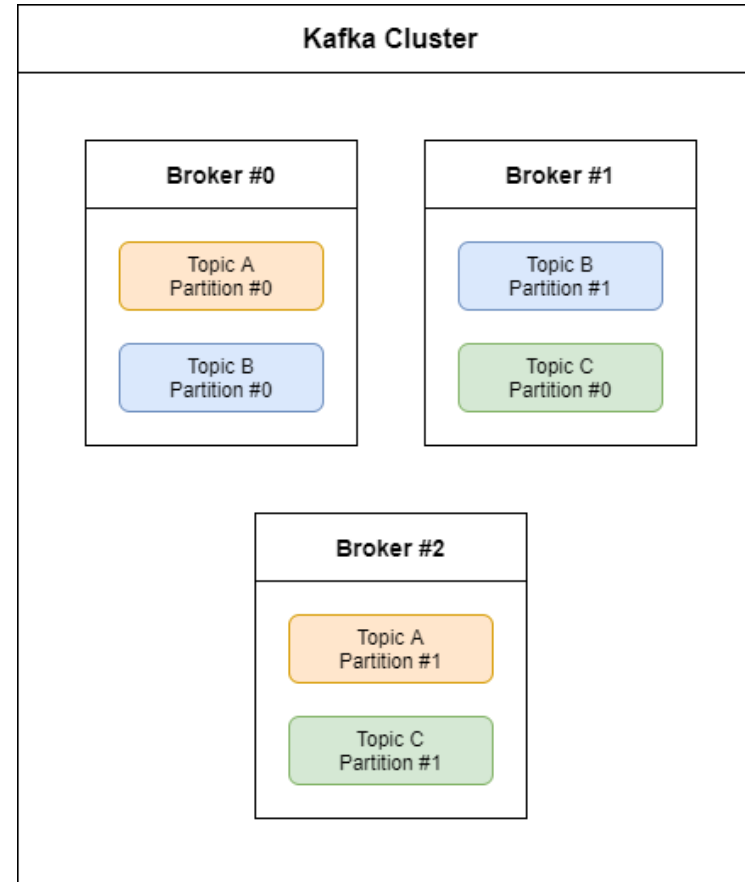
Brokers

- ▶ Brokers can be
 - ▶ Physical servers
 - ▶ Containers
 - ▶ Cloud instances
- ▶ Manage partitions
- ▶ Communicate with other brokers
- ▶ They are simple
 - ▶ Understand easily
 - ▶ Scale easily
 - ▶ Extend easily



Kafka cluster

- ▶ Kafka is a distributed system
- ▶ Good scalability
- ▶ Many brokers form a Kafka cluster
- ▶ From outside, we see it as „one” Kafka




Kafka records

- ▶ Events represented in Kafka are called records
- ▶ Structure of a record
 - ▶ Key
 - ▶ Value
 - ▶ Headers
 - ▶ Timestamp

Record
Headers
Key
Value
Timestamp

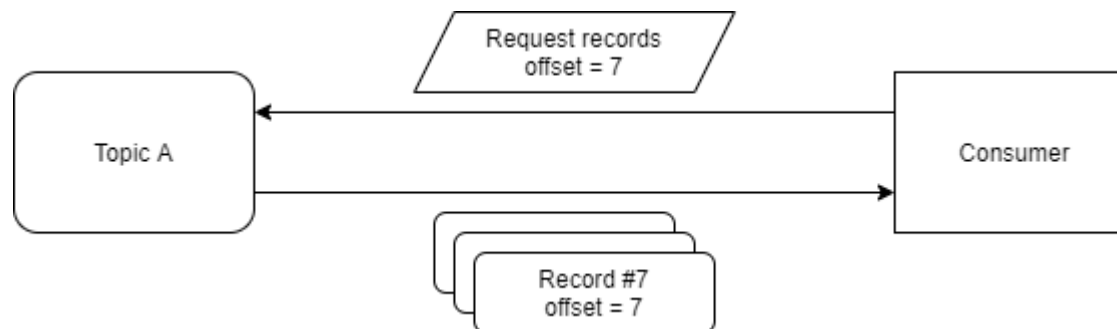
Producers

- ▶ Applications that we write
- ▶ Send records to topics (partitions)
- ▶ Producer API
 - ▶ Connection pooling
 - ▶ Networking
 - ▶ Partitioning

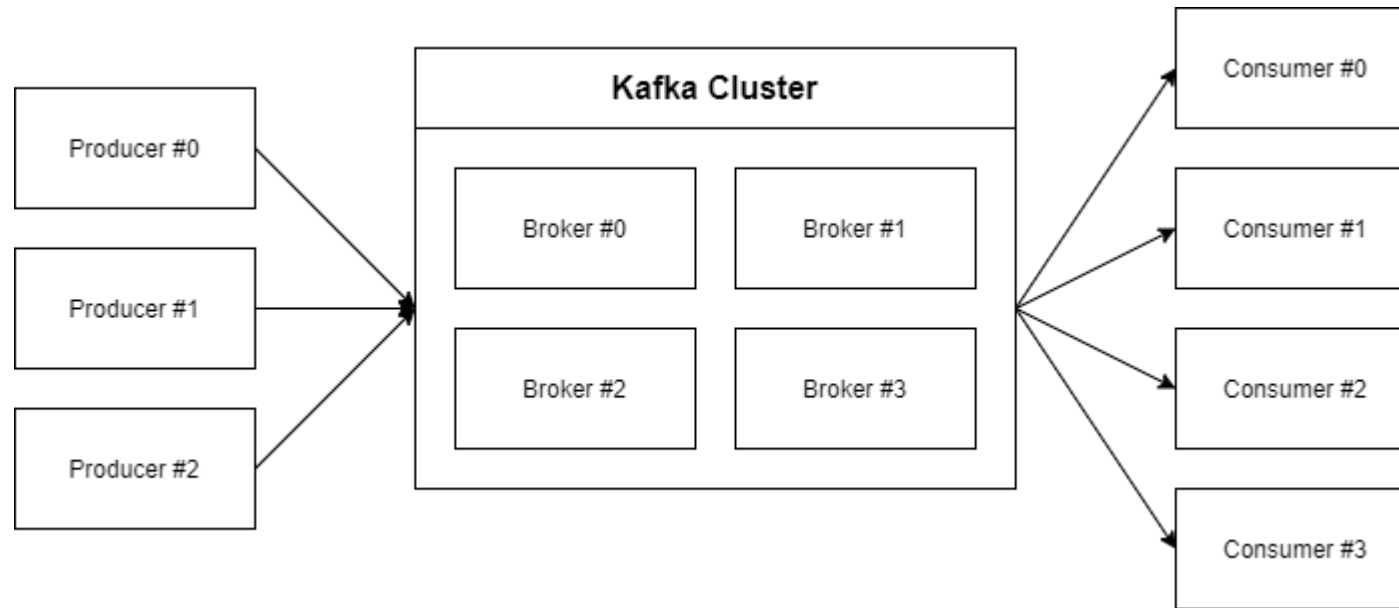


Consumers

- ▶ Applications that we write
- ▶ Read (pull) records from topics
 - ▶ Read by offset
- ▶ Consumer API
 - ▶ Connection pooling
 - ▶ Networking
- ▶ A record doesn't get destroyed after read




Fundamental parts of Kafka



Let's take a closer look!

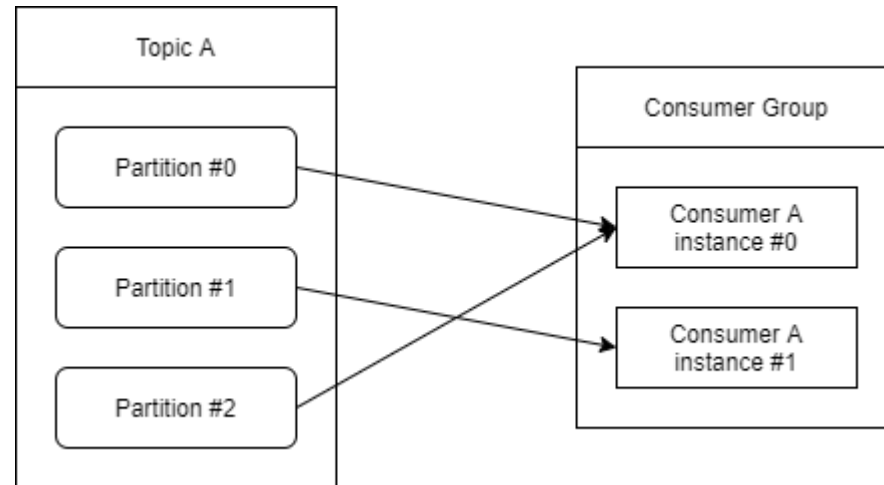
We produce to partitions

- ▶ Records are sent to partitions
- ▶ Purpose
 - ▶ Load balancing
 - ▶ Semantic partitioning
- ▶ Record without key
 - ▶ Round-robin strategy
- ▶ Record with key
 - ▶ $\text{hash}(\text{key}) \% \text{num_of_partitions}$
 - ▶ Records with same key to same partition
- ▶ Custom partitioner logic



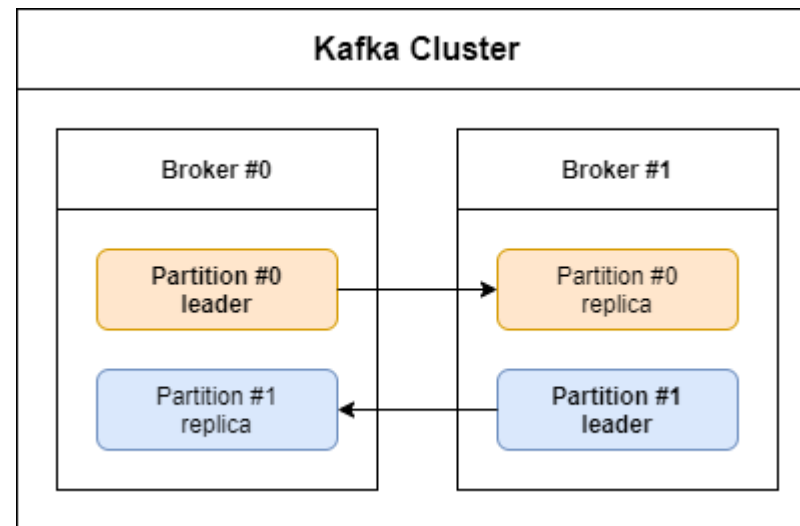
Consumer groups

- ▶ A consumer reads from partition(s)
- ▶ We might want to scale a consumer application
- ▶ Every consumer is part of a consumer group
- ▶ Kafka handles rebalancing automatically



Replication

- ▶ Brokers store records (in partitions, segments)
- ▶ Brokers might go down
- ▶ Replication factor
 - ▶ 1 leader partition
 - ▶ N-1 follower partition
- ▶ Synchronization is automatic
- ▶ Produce only to leaders
- ▶ We can configure to read from followers

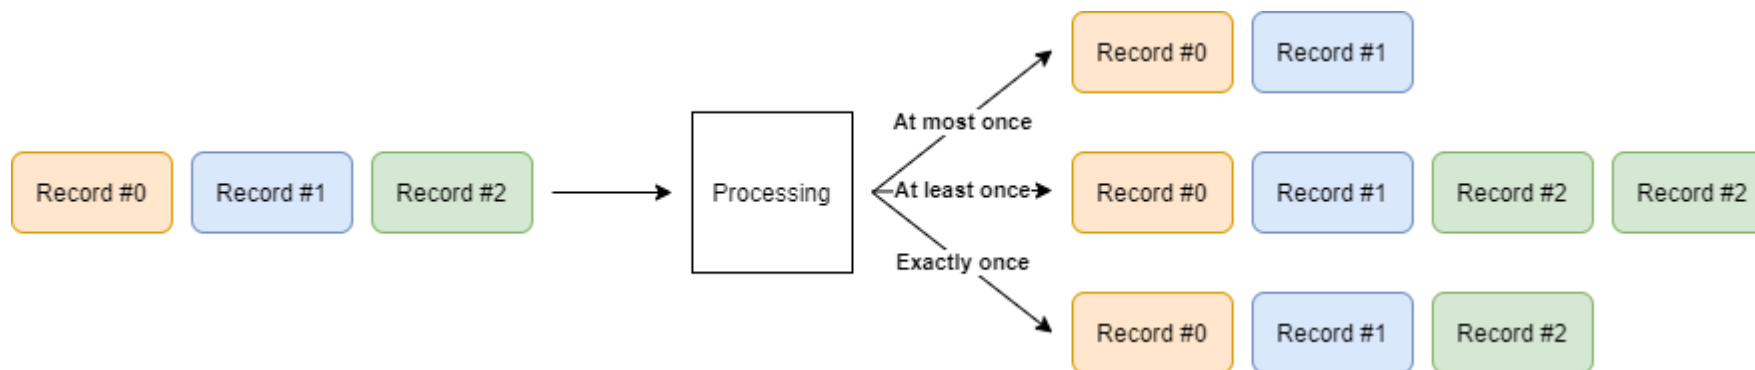


Producer guarantees

- ▶ We can wait for an acknowledgement from the broker(s)
 - ▶ Is the record successfully written to disk?
- ▶ This is a producer configuration
- ▶ We have 3 options
 - ▶ **NONE**
 - ▶ Fast, but might lose data
 - ▶ **LEADER**
 - ▶ Default, data is written for sure
 - ▶ **ALL**
 - ▶ Slowest option

Delivery guarantees

- ▶ Message processing guarantee categories
 - ▶ At most once
 - ▶ At least once
 - ▶ Exactly once
- ▶ Kafka provides a transactional API to reach exactly once
- ▶ Configurable for producer and consumer



Retention policy

- ▶ For how long a record should be stored?
- ▶ Decision factor
 - ▶ Business
 - ▶ Cost
- ▶ Default: 1 week
- ▶ Configurable
 - ▶ Globally
 - ▶ Per topic
- ▶ Data purge happens per segment



Security

- ▶ Kafka supports data encryption in transit
- ▶ Supports authentication and authorization
- ▶ Data stored in disk by Kafka is not encrypted
 - ▶ We can encrypt the disk itself
 - ▶ We can write a wrapper to encrypt/decrypt a record's value



The Kafka ecosystem

The background of the slide features abstract geometric shapes in various shades of teal and grey, primarily concentrated on the right side, creating a modern and layered visual effect.

ZooKeeper

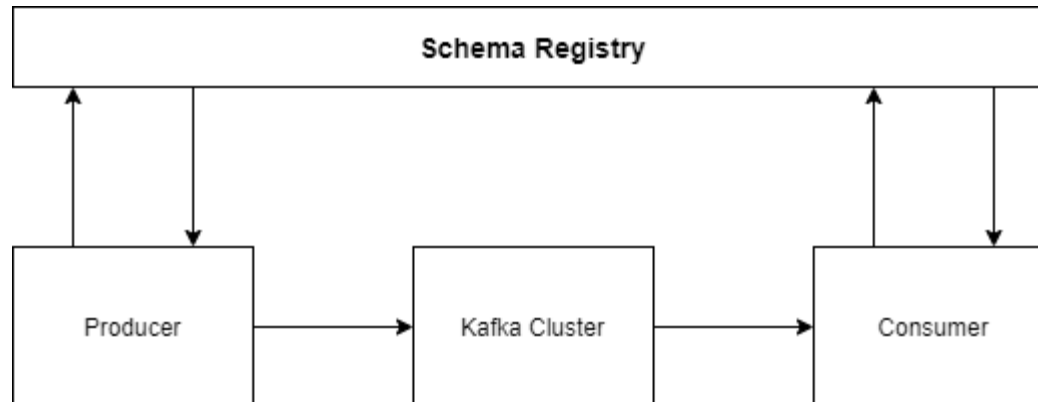
- ▶ Provides multiple features for distributed applications
 - ▶ Distributed configuration management
 - ▶ Election
 - ▶ Key-value store
- ▶ Kafka uses ZooKeeper
 - ▶ Leader election
 - ▶ Topic configuration
 - ▶ Access control
- ▶ KIP-500, Apache Kafka 2.8
 - ▶ We can leave ZooKeeper
 - ▶ Not production ready yet



APACHE
ZooKeeper[™]

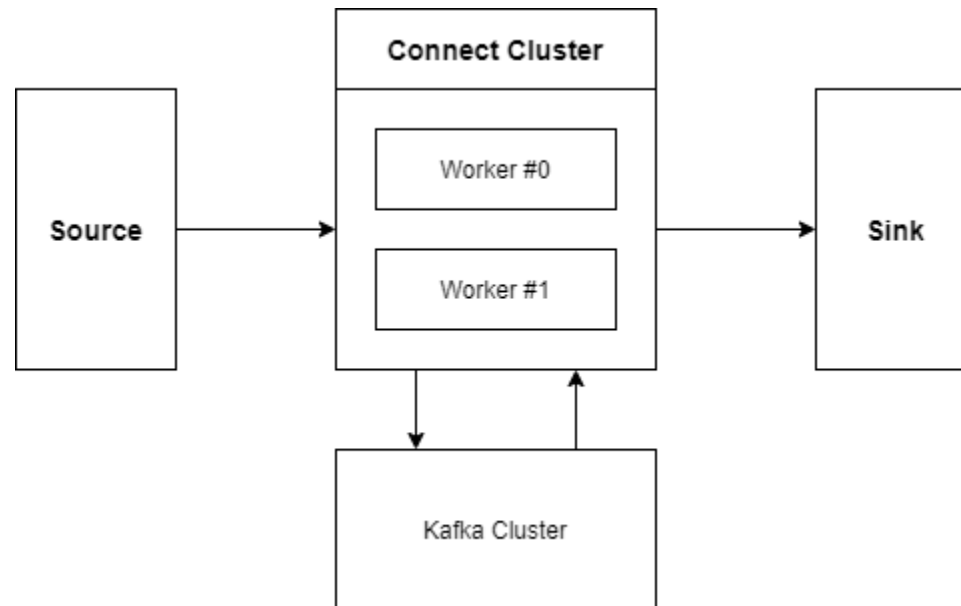
Schema Registry

- ▶ Schema of or records might change by time
- ▶ A schema registry is a database of schemas
- ▶ Supported formats: AVRO, JSON schema, Protobuf, custom
- ▶ Producers and consumers validate the records
 - ▶ They call the schema registry
 - ▶ Check the compatibility



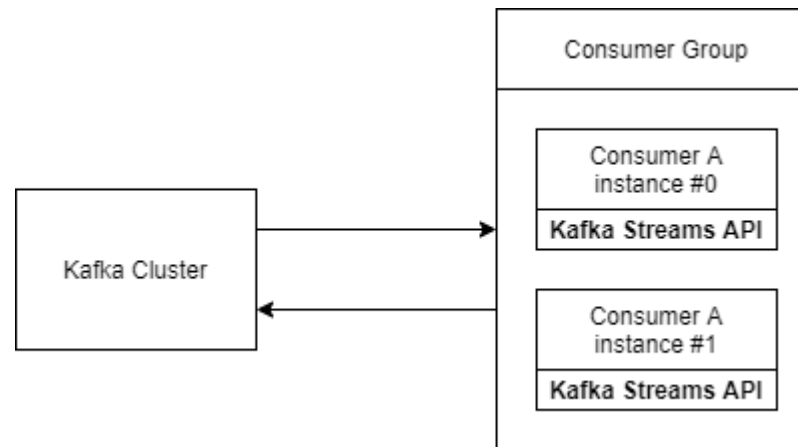
Kafka Connect

- ▶ We might want to transfer data
 - ▶ From Kafka to another system
 - ▶ Into Kafka from another system
- ▶ Many existing connectors
- ▶ Easy to configure




Kafka Streams

- ▶ The consumer API is very simple
- ▶ Kafka Streams provides more complex operations
 - ▶ Filtering
 - ▶ Mapping
 - ▶ Aggregating
- ▶ Fluent functional API
- ▶ Output is stored in Kafka
- ▶ Scalable, fault-tolerant



ksqlDB

- ▶ Another stream processing option
- ▶ Without programming
- ▶ Stream processing in an SQL format
- ▶ Communicate with ksqlDB
 - ▶ CLI
 - ▶ GUI



REST Proxy

- ▶ Java is the native language for Kafka
- ▶ Lot of other supported languages
- ▶ We can use the REST Proxy instead of libraries
 - ▶ HTTP calls for producing and consuming
 - ▶ Useful if our chosen language is not supported



Sources and result

► Sources

- [Confluent - Apache Kafka® Tutorials | Kafka 101](#)
- [Confluent - Course | Apache Kafka® Fundamentals](#)
- <https://kafka.apache.org/>

► Result

- <https://github.com/veresdavid/apache-kafka-basics>



Thank you for your
attention!